treeIDW

treeIDW is a Python library designed for performing Inverse Distance Weighting (IDW) interpolation with an efficient KD-tree approach. It offers a user-friendly interface for newcomers while providing advanced features and optimizations that will appeal to experts in numerical methods and spatial data analysis.

Introduction

Imagine you have data known at certain points (boundary nodes) and you need to estimate values at other points (internal nodes). **treeIDW** simplifies this task by automatically selecting the most relevant boundary nodes using a KD-tree, ensuring that only nodes with significant contributions are considered. This not only makes the interpolation more accurate but also speeds up the computation considerably. Whether you're a researcher, engineer, or someone exploring spatial data, **treeIDW** offers a robust solution that adapts to your level of expertise.

For non-experts, the library hides complex mathematical operations behind a simple interface, while experts will appreciate the fine-tuned control over parameters such as weight thresholds, bisection tolerances, and parallel processing options.

Installation

Since **treeIDW** is not yet available on PyPI, you can install it locally as follows:

```
git clone https://github.com/Dorian210/treeIDW
cd treeIDW
pip install -e .
```

Dependencies

Ensure that you have the following dependencies installed:

- numpy
- numba
- scipy

Main Modules

- **treeIDW.treeIDW**
  Implements the core IDW interpolation method using a KD-tree to select the most relevant boundary nodes. This module efficiently ignores nodes with negligible impact, improving both accuracy and performance.

- **treeIDW.helper_functions**
  Contains optimized functions (leveraging `numba`) for computing IDW weights. These include both vectorized and parallelized implementations, which are ideal for handling large datasets.

- **treeIDW.weight_function**
  Defines the specific IDW weight calculation function. By default, this function computes the weight as the inverse of the squared distance, but it can be adapted to suit specific needs.

Examples

Several example scripts demonstrating the usage of **treeIDW** are available in the `examples/` directory. These include:

- **Graphical Demonstration:** A visualization of a square domain with a rotating vector field. The field is propagated to multiple internal points and plotted for intuitive understanding.
- **Large-Scale Computation:** A more computationally intensive example where a scalar field is propagated from 1,000 boundary nodes to 1,000,000 internal nodes, showcasing the efficiency of the KD-tree selection.
- **Logo Generation:** A unique example illustrating the process of creating the treeIDW logo. The logo consists of the letters "IDW" represented as a vector field, which is then propagated onto a 2D meshgrid to generate the final design.

Documentation

Complete API documentation is available in the `docs/` directory of the project or through the online documentation portal.

Contributions

Currently, I am not actively reviewing contributions. However, if you encounter any issues or have suggestions, please feel free to open an issue on the repository.

License

This project is licensed under the CeCILL License.

- View Source

treeIDW.weight_function

```
@nb.njit(cache=True)
def compute_weight(dist_squared: float) -> float:
```

Compute the IDW weight. Can be changed but is usually `1/dist_squared` .

Parameters

- **dist_squared** (float): Squared distance between the interior node and the boundary node.

Returns

- **weight** (float): IDW weight.

treeIDW.weight_function

```
@nb.njit(cache=True)
def compute_weight(dist_squared: float) -> float:
```

treeIDW.treeIDW

```
def treeIDW(
 boundary_nodes: numpy.ndarray[numpy.floating],
 boundary_field: numpy.ndarray[numpy.floating],
 internal_nodes: numpy.ndarray[numpy.floating],
 neglectible_treshold: float = 0.2,
 bisect_rtol: float = 0.001,
 parallel: bool = False
) -> numpy.ndarray[numpy.floating]:
```

Performs IDW interpolation using a KD-tree to select relevant boundary nodes for each internal node. Only boundary nodes with significant weights are included in the interpolation.

Parameters

- **boundary_nodes** (np.ndarray[np.floating]): Boundary nodes where the data is known. Must be of shape (`n_interp` , `space_dim`).
- **boundary_field** (np.ndarray[np.floating]): Known data. Must be of shape (`n_interp` , `field_dim`).
- **internal_nodes** (np.ndarray[np.floating]): Internal nodes where the interpolator is evaluated. Must be of shape (`n_eval` , `space_dim`).
- **neglectible_treshold** (float, optional): Relative weight threshold below which boundary nodes are ignored, by default 0.2
- **bisect_rtol** (float, optional): Relative tolerance for bisection convergence, by default 1e-3
- **parallel** (bool, optional): Whether to use parallel implementation, by default False

Returns

- **internal_field** (np.ndarray[np.floating]): Interpolated data. Should be of shape (`n_eval` , `field_dim`).

treeIDW.helper_functions

```
@nb.njit(cache=True)
def inv_dist_weight(
 boundary_nodes: numpy.ndarray[numpy.floating],
 boundary_field: numpy.ndarray[numpy.floating],
 internal_nodes: numpy.ndarray[numpy.floating],
 relevant_nodes_inds_flat: numpy.ndarray[numpy.integer],
 relevant_nodes_inds_sizes: numpy.ndarray[numpy.integer]
) -> numpy.ndarray[numpy.floating]:
```

Performs the Inverse Distance Weighting interpolation method using only provided boundary nodes indices in the weighted sum. This function is designed to be used after selecting the relevant nodes through a KD-tree method.

Parameters

- **boundary_nodes** (np.ndarray[np.floating]): Boundary nodes where the data is known. Must be of shape (`n_interp` , `space_dim`).
- **boundary_field** (np.ndarray[np.floating]): Known data. Must be of shape (`n_interp` , `field_dim`).
- **internal_nodes** (np.ndarray[np.floating]): Internal nodes where the interpolator is evaluated. Must be of shape (`n_eval` , `space_dim`).
- **relevant_nodes_inds_flat** (np.ndarray[np.integer]): Boundary nodes indices that weight in the IDW interpolator for each internal node. Must contain `n_eval` concatenated lists of indices between 0 and `n_interp` excluded.
- **relevant_nodes_inds_sizes** (np.ndarray[np.integer]): Sizes of each of the `n_eval` lists concatenated in `relevant_nodes_inds_flat` .

Returns

- **internal_field** (np.ndarray[np.floating]): Interpolated data. Should be of shape (`n_eval` , `field_dim`).

```
@nb.njit(parallel=True, cache=True)
def inv_dist_weight_parallel(
 boundary_nodes: numpy.ndarray[numpy.floating],
 boundary_field: numpy.ndarray[numpy.floating],
 internal_nodes: numpy.ndarray[numpy.floating],
 relevant_nodes_inds_flat: numpy.ndarray[numpy.integer],
 relevant_nodes_inds_sizes: numpy.ndarray[numpy.integer]
) -> numpy.ndarray[numpy.floating]:
```

Performs the Inverse Distance Weighting interpolation method using only provided boundary nodes indices in the weighted sum. This function is designed to be used after selecting the relevant nodes through a KD-tree method.

Parameters

- **boundary_nodes** (np.ndarray[np.floating]): Boundary nodes where the data is known. Must be of shape (`n_interp` , `space_dim`).
- **boundary_field** (np.ndarray[np.floating]): Known data. Must be of shape (`n_interp` , `field_dim`).
- **internal_nodes** (np.ndarray[np.floating]): Internal nodes where the interpolator is evaluated. Must be of shape (`n_eval` , `space_dim`).
- **relevant_nodes_inds_flat** (np.ndarray[np.integer]): Boundary nodes indices that weight in the IDW interpolator for each internal node. Must contain `n_eval` concatenated lists of indices between 0 and `n_interp` excluded.
- **relevant_nodes_inds_sizes** (np.ndarray[np.integer]): Sizes of each of the `n_eval` lists concatenated in `relevant_nodes_inds_flat` .

Returns

- **internal_field** (np.ndarray[np.floating]): Interpolated data. Should be of shape (`n_eval` , `field_dim`).

```
@nb.njit(cache=True)
def bisect_weight_elem(
 dist_squared_a: float,
 dist_squared_b: float,
 weight_treshold: float,
 rtol: float
) -> float:
```

Compute the squared distance that correspond to a specific weight threshold using a bisection approach.

Parameters

- **dist_squared_a** (float): Left boundary of the interval containning the looked for weight treshold preimage.
- **dist_squared_b** (float): Right boundary of the interval containning the looked for weight treshold preimage.
- **weight_treshold** (float): Image of the looked for squared distance by the weight function of the IDW algorithm.
- **rtol** (float, optional): The relative tolerance is used to compute the necessary number of bisection iterations, by default 1e-3

Returns

- **dist_squared_c** (float): Weight treshold preimage found through bisection.

 bisect_weight = <numba._GUFunc 'bisect_weight'>

Vectorized version of bisect_weight_elem that finds squared distances corresponding to weight thresholds.

Parameters

- **dist_squared_a** (float): Left boundary of interval containing weight threshold preimage
- **dist_squared_b** (float): Right boundary of interval containing weight threshold preimage
- **weight_treshold** (float): Target weight threshold value
- **rtol** (float): Relative tolerance for bisection convergence
- **out** (ndarray): Output array for storing computed squared distance

 bisect_weight_parallel = <ufunc 'bisect_weight_parallel'>

Vectorized version of bisect_weight_elem that finds squared distances corresponding to weight thresholds. Parallelized version of the function.

Parameters

- **dist_squared_a** (float): Left boundary of interval containing weight threshold preimage
- **dist_squared_b** (float): Right boundary of interval containing weight threshold preimage
- **weight_treshold** (float): Target weight threshold value
- **rtol** (float): Relative tolerance for bisection convergence
- **out** (ndarray): Output array for storing computed squared distance

 compute_weight_vectorized = <numba._GUFunc 'compute_weight_vectorized'>

Vectorized version of compute_weight function for IDW calculations.

Parameters

- **dist_squared** (float): Squared distance between interior and boundary nodes.
- **out** (ndarray): Output array containing the computed IDW weights.

 compute_weight_vectorized_parallel = <ufunc 'compute_weight_vectorized_parallel'>

Vectorized version of compute_weight function for IDW calculations. Parallelized version of the function.

Parameters

- **dist_squared** (float): Squared distance between interior and boundary nodes.
- **out** (ndarray): Output array containing the computed IDW weights.